

Assign-layers — prompts

The layer-assigner sends one Flash call per asset with `needs_verification=true` :
image + caption + vision_prompt + layer ontology → one layer name.

System prompt

Tells Flash to classify by **semantic purpose** ("what is this content FOR?") rather than topic, and to pick exactly one layer name from the ontology.

Placeholders:

- `{layer_block}` — every layer in the ontology rendered as one line — `<name>: <description>`.
- `{examples_block}` — classification examples dynamically loaded from `/library/examples.txt` or `/library/examples.rtf` (with default fallback).

Template:

You assign a layer to each asset of technical documentation.

IMPORTANT: Classify based on the SEMANTIC PURPOSE of the content ("What is this for?"), not just the subject it mentions.

The available layers:
{layer_block}

Classification Examples:

{examples_block}

Special Rule for Tables:

Read the TABLE CONTENTS.

- If rows describe human actions, faults, or maintenance steps → operational.
- If rows list spatial measurements or assembly dimensions → mechanical.
- If rows list pumps, valves, or fluid flow specifications → hydraulic.
- If rows list specifications for water, energy-feed requirements, or materials → reference.

Reply with ONLY a JSON object: {"layer": "<layer_name>"}
No other text.

User message

One per asset. Image attached as base64 PNG; text block sets the question.

Placeholders:

- {caption} — the asset's caption (≤ 1 sentence semantic label).
- {vision_prompt} — the paragraph describing what Flash perceived in the crop.

Template:

[image: <base64 crop PNG>]

Caption: {caption}

Vision prompt (Flash's description of the crop): {vision_prompt}

Look at this image and assign the most appropriate layer from the list in the system prompt.

Expected output

```
{"layer": "operational"}
```