

Caption — prompts

The captioner runs one Gemini API call per HITL-approved region (Gemini 3.5 Flash, no tools): reads the crop and emits JSON with `vision_prompt` and `caption`, written back to `regions.yaml`. Two prompt surfaces reach the model.

System prompt

Frames the reading task, embeds the captioning rules verbatim, and closes with the JSON output spec.

Placeholders:

- `{captioning_rules}` — the `## B` section of `rag/rules.md` (captioning + vision prompts), metadata anchors stripped.

Template:

```
You are reading a single image crop. Apply the captioning rules below:
they define what to output, what to ignore, and what each field should
contain.
```

```
— RULES (verbatim from rag/rules.md § B) —
```

```
{captioning_rules}
```

```
— OUTPUT FORMAT —
```

```
Reply with ONLY a JSON object, no other text:
{"vision_prompt": "<paragraph>", "caption": "<short label>"}
```

User message per region

One per region under captioning. Crop attached as base64 image (PNG or JPEG inferred from suffix); text block is fixed.

No placeholders.

Template:

```
[image: <base64 crop>]
```

Read the image. Write vision_prompt first, then caption.

Expected output

```
{"vision_prompt": "Paragraph describing the semantic meaning of the  
main subject...", "caption": "Short semantic label"}
```